

CENG 331

Computer Organization

Fall '2015-2016

Pre-lab Exam 4

Prelab due date: 04 January 2016, Monday, 17:00

Exam date: 04 January 2016, Monday, 17:40

1 Objectives

This assignment deals with optimizing memory intensive code. Image processing offers many examples of functions that can benefit from optimization. In this lab, we will consider two image processing operations: **rotate**, which rotates an image clockwise by 90° , and **smooth by masking**, which “smooths” or “blurs” an image.

For this lab, we will consider an image to be represented as a two-dimensional matrix M , where $M_{i,j}$ denotes the value of (i,j) th pixel of M . Pixel values are triples of red, green, and blue (RGB) values. We will only consider square images. Let N denote the number of rows (or columns) of an image. Rows and columns are numbered, in C-style, from 0 to $N - 1$.

Given this representation, the **rotate** operation can be implemented quite simply as the combination of the following two matrix operations:

- *Transpose*: For each (i,j) pair, $M_{i,j}$ and $M_{j,i}$ are interchanged.
- *Exchange columns*: Column i is exchanged with column $N - 1 - i$.

This combination is illustrated in Figure 1.

The **smooth by masking** operation is implemented by replacing every pixel value with the weighted average of all the pixels around it (in a maximum of 3×3 window centered at that pixel). Consider Figure 2. The values of pixels $M2[1][1]$ and $M2[N-1][N-1]$ are given below:

$$M2[1][1] = \frac{M1[0][0] \times 1 + M1[0][1] \times 2 + M1[0][2] \times 1 + M1[1][0] \times 2 + M1[1][1] \times 4 + M1[1][2] \times 2 + M1[2][0] \times 1 + M1[2][1] \times 2 + M1[2][2] \times 1}{16}$$

$$M2[N-1][N-1] = \frac{M1[N-2][N-2] \times 1 + M1[N-2][N-1] \times 2 + M1[N-1][N-2] \times 2 + M1[N-1][N-1] \times 4}{8}$$

Keywords: *Performance, Optimization, Loop Unrolling, Loop Breaking*




Figure 1: Rotation of an image by 90° clockwise




Figure 2: Smoothing an image by masking

2 Specifications

Start by copying `perflab-handout.tar` to a protected directory in which you plan to do your work. Then give the command: `tar xvf perflab-handout.tar`. This will cause a number of files to be unpacked into the directory. The only file you will be modifying and handing in is `kernels.c`. The `driver.c` program is a driver program that allows you to evaluate the performance of your solutions. Use the command `make driver` to generate the driver code and run it with the command `./driver`.

Looking at the file `kernels.c` you'll notice a C structure `student` into which you should insert the requested identifying information about you. **Do this right away so you don't forget.**

3 Implementation Overview

Data Structures

The core data structure deals with image representation. A `pixel` is a struct as shown below:

```
typedef struct {
    unsigned short red;    /* R value */
    unsigned short green; /* G value */
    unsigned short blue;  /* B value */
} pixel;
```

As can be seen, RGB values have 16-bit representations ("16-bit color"). An image `I` is represented as a one-dimensional array of `pixels`, where the (i, j) th pixel is `I[RIDX(i, j, n)]`. Here `n` is the dimension of the image matrix, and `RIDX` is a macro defined as follows:

```
#define RIDX(i, j, n) ((i)*(n)+(j))
```

See the file `defs.h` for this code.

Rotate

The following C function computes the result of rotating the source image `src` by 90° and stores the result in destination image `dst`. `dim` is the dimension of the image.

```
void naive_rotate(int dim, pixel *src, pixel *dst) {
    int i, j;

    for(i=0; i < dim; i++)
        for(j=0; j < dim; j++)
            dst[RIDX(j, dim-1-i, dim)] = src[RIDX(i, j, dim)];

    return;
}
```

The above code scans the columns of the source image matrix, copying to the rows of the destination image matrix. Your task is to rewrite this code to make it run as fast as possible using techniques like code motion, loop unrolling and blocking.

See the file `kernels.c` for this code.

Test case	1	2	3	4	5	
Method N	64	128	256	512	1024	Geom. Mean
Naive rotate (CPE)	2.1	2.1	2.2	2.9	4.5	
Optimized rotate (CPE)	2.8	4.1	6.4	10.3	11.3	
Speedup (naive/opt)	1.3	1.9	2.9	3.5	2.5	2.3
Method N	32	64	128	256	512	Geom. Mean
Naive smooth (CPE)	19.9	21.6	22.9	24.3	22.9	
Optimized smooth (CPE)	79.5	79.6	80.0	80.2	80.1	
Speedup (naive/opt)	4.0	3.7	3.5	3.3	3.3	3.5

Table 1: CPEs and Ratios for Optimized vs. Naive Implementations

Smooth by masking

The smoothing function takes as input a source image **src** and returns the smoothed result in the destination image **dst**. Here is part of an implementation:

```
void naive_smooth(int dim, pixel *src, pixel *dst) {
    int i, j;

    for(i=0; i < dim; i++)
        for(j=0; j < dim; j++)
            dst[RIDX(i,j,dim)] = avg(dim, i, j, src); /* Smooth the (i,j)th pixel */

    return;
}
```

The function **avg** returns the weighted average of all the pixels around the (i,j)th pixel. Your task is to optimize **smooth** (and **avg**) to run as fast as possible. (*Note:* The function **avg** is a local function and you can get rid of it altogether to implement **smooth** in some other way.)

This code (and an implementation of **avg**) is in the file **kernels.c**.

Performance measures

Our main performance measure is *CPE* or *Cycles per Element*. If a function takes C cycles to run for an image of size $N \times N$, the CPE value is C/N^2 . Table 1 summarizes the performance of the naive implementations shown above and compares it against an optimized implementation. Performance is shown for 5 different values of N . All measurements were made on the Intel(R) Core(TM) i7-4770S CPU @ 3.1GHz.

The ratios (speedups) of the optimized implementation over the naive one will constitute a *score* of your implementation. To summarize the overall effect over different values of N , we will compute the *geometric mean* of the results for these 5 values. That is, if the measured speedups for $N = \{32, 64, 128, 256, 512\}$ are R_{32} , R_{64} , R_{128} , R_{256} , and R_{512} then we compute the overall performance as

$$R = \sqrt[5]{R_{32} \times R_{64} \times R_{128} \times R_{256} \times R_{512}}$$

Assumptions

To make life easier, you can assume that N is a multiple of 32. Your code must run correctly for all such values of N , but we will measure its performance only for the 5 values shown in Table 1.

4 Infrastructure

We have provided support code to help you test the correctness of your implementations and measure their performance. This section describes how to use this infrastructure. The exact details of each part of the assignment is described in the following section.

Note: The only source file you will be modifying is `kernels.c`.

Versioning

You will be writing many versions of the `rotate` and `smooth` routines. To help you compare the performance of all the different versions you’ve written, we provide a way of “registering” functions.

For example, the file `kernels.c` that we have provided you contains the following function:

```
void register_rotate_functions() {
    add_rotate_function(&rotate, rotate_descr);
}
```

This function contains one or more calls to `add_rotate_function`. In the above example, `add_rotate_function` registers the function `rotate` along with a string `rotate_descr` which is an ASCII description of what the function does. See the file `kernels.c` to see how to create the string descriptions. This string can be at most 256 characters long.

A similar function for your smooth by masking kernels is provided in the file `kernels.c`.

Driver

The source code you will write will be linked with object code that we supply into a `driver` binary. To create this binary, you will need to execute the command

```
unix> make driver
```

You will need to re-make driver each time you change the code in `kernels.c`. To test your implementations, you can then run the command:

```
unix> ./driver
```

The `driver` can be run in four different modes:

- *Default mode*, in which all versions of your implementation are run.
- *Autograder mode*, in which only the `rotate()` and `smooth()` functions are run. This is the mode we will run in when we use the driver to grade your handin.
- *File mode*, in which only versions that are mentioned in an input file are run.
- *Dump mode*, in which a one-line description of each version is dumped to a text file. You can then edit this text file to keep only those versions that you’d like to test using the *file mode*. You can specify whether to quit after dumping the file or if your implementations are to be run.

If run without any arguments, `driver` will run all of your versions (*default mode*). Other modes and options can be specified by command-line arguments to `driver`, as listed below:

`-g` : Run only `rotate()` and `smooth()` functions (*autograder mode*).

`-f <funcfile>` : Execute only those versions specified in `<funcfile>` (*file mode*).

- d <dumpfile> : Dump the names of all versions to a dump file called <dumpfile>, *one line* to a version (*dump mode*).
- q : Quit after dumping version names to a dump file. To be used in tandem with -d. For example, to quit immediately after printing the dump file, type `./driver -qd dumpfile`.
- h : Print the command line usage.

Student Information

Important: Before you start, you should fill in the struct in `kernels.c` with information about you (student name, student id).

5 Assignment Details

Optimizing Rotate (50 points)

In this part, you will optimize `rotate` to achieve as low a CPE as possible. You should compile `driver` and then run it with the appropriate arguments to test your implementations.

For example, running `driver` with the supplied naive version (for `rotate`) generates the output shown below:

```
unix> ./driver
Student Name: Alperen Dalkıran
Student Id: eXXXXXXX

Rotate: Version = naive_rotate: Naive baseline implementation:
Dim          64      128      256      512      1024      Mean
Your CPEs    2.6      3.6      6.1      10.0      11.2
Baseline CPEs 2.8      4.1      6.4      10.3      11.3
Speedup      1.0      1.1      1.0      1.0      1.0      1.0
```

Optimizing Smooth by Masking (50 points)

In this part, you will optimize `smooth by masking` to achieve as low a CPE as possible.

For example, running `driver` with the supplied naive version (for `smooth by masking`) generates the output shown below:

```
unix> ./driver

Smooth: Version = naive_smooth: Naive baseline implementation:
Dim          32      64      128      256      512      Mean
Your CPEs    78.7    79.1    79.5    79.5    79.9
Baseline CPEs 79.5    79.6    80.0    80.2    80.1
Speedup      1.0      1.0      1.0      1.0      1.0      1.0
```

Some advice. Look at the assembly code generated for the `rotate` and `smooth by masking`. Focus on optimizing the inner loop (the code that gets repeatedly executed in a loop) using the optimization tricks covered in class. The `smooth by masking` is more compute-intensive and less memory-sensitive than the `rotate` function, so the optimizations are of somewhat different flavors.

Coding Rules

You may write any code you want, as long as it satisfies the following:

- It must be in ANSI C. You may not use any embedded assembly language statements.
- It must not interfere with the time measurement mechanism. You will also be penalized if your code prints any extraneous information.

You can only modify code in `kernels.c`. You are allowed to define macros, additional global variables, and other procedures in these files.

Evaluation

Your solutions for `rotate` and `smooth` will each count for 50% of your grade. The score for each will be based on the following:

- **Correctness:** You will get NO CREDIT for buggy code that causes the driver to complain! This includes code that correctly operates on the test sizes, but incorrectly on image matrices of other sizes. As mentioned earlier, you may assume that the image dimension is a multiple of 32.
- **Speed-up:** You will get 20 for each implementation of `rotate` and `smooth` if they are correct and achieve mean speed-up threshold of 1.4 and 3.0 respectively for rotate and smooth by masking. One who achieve the biggest speed-up will get 50 points for each implementation. Other grades will be scaled between 20 and 50 according to your speed-up for each function. You will not get partial credit for a correct implementation that does below the threshold.
IMPORTANT: Same rules also valid for in-lab exam.

6 Regulations

1. **Programming Language:** You must write program in ANSI C.
2. **Submission:** Submission will be done via COW. Submit a single c source file named `kernels.c` which will be modified version of supplied `bits.c` source file.
3. **Late Submission:** Late submission is not allowed.
4. **Cheating: We have zero tolerance policy for cheating.** People involved in cheating will be punished according to the university regulations.
5. **Newsgroup:** You must follow the newsgroup (news.ceng.metu.edu.tr) for discussions and possible updates on a daily basis.